

CSCI596 Assignment 6—Hybrid MPI+OpenMP+CUDA Programming

Youzhi Qu

CSCI-596 assignment4

Prof. Aiichiro Nakano

Part 1:

pdf1.cu code:

```
/*-----
Program pdf0.c computes a pair distribution function for n atoms
given the 3D coordinates of the atoms.
-----*/

#include <stdio.h>
#include <math.h>
#include <time.h>
#include <stdlib.h>

#define NHBIN 2000 // Histogram size

float al[3]; // Simulation box lengths
int n; // Number of atoms
float *r; // Atomic position array
FILE *fp;

__constant__ float DALTH[3];
__constant__ int DN;
__constant__ float DDRH;

//float SignR(float v,float x) {if (x > 0) return v; else return -v;}
__device__ float d_SignR(float v,float x) {if (x > 0) return v; else return -v;}

__global__ void gpu_histogram_kernel(float *r,float *nhis) {
    int i,j,a,ih;
    float rij,dr;

    int iBlockBegin = (DN/gridDim.x)*blockIdx.x;
    int iBlockEnd = min((DN/gridDim.x)*(blockIdx.x+1),DN);

    int jBlockBegin = (DN/gridDim.y)*blockIdx.y;
    int jBlockEnd = min((DN/gridDim.y)*(blockIdx.y+1),DN);

    for (i=iBlockBegin+threadIdx.x; i<iBlockEnd; i+=blockDim.x) {
        for (j=jBlockBegin+threadIdx.y; j<jBlockEnd; j+=blockDim.y) {
```

```

        if (i<j) {
        // Process (i,j) atom pair
            rij = 0.0;
            for (a=0; a<3; a++) {
                dr = r[3*i+a]-r[3*j+a];
                /* Periodic boundary condition */
                dr=dr-d_SignR(DALTH[a],dr-DALTH[a])-
d_SignR(DALTH[a],dr+DALTH[a]);
                rij += dr*dr;
            }
            rij = sqrt(rij); /* Pair distance */
            ih = rij/DDRH;
            //nhis[ih] += 1.0;
            atomicAdd(&nhis[ih],1.0);
        } // end if i<j
    } // end for j
} // end for i
}
/*-----*/
void histogram() {
/*-----*/
Constructs a histogram NHIS for atomic-pair distribution.
-----*/

    float alth[3];
    float* nhis; // Histogram array
    float rhmax,drh,density,gr;
    int a,ih;

    float* dev_r; // Atomic positions
    float* dev_nhis; // Histogram

    /* Half the simulation box size */
    for (a=0; a<3; a++) alth[a] = 0.5*al[a];
    /* Max. pair distance RHMAX & histogram bin size DRH */
    rhmax = sqrt(alth[0]*alth[0]+alth[1]*alth[1]+alth[2]*alth[2]);
    drh = rhmax/NHBIN; // Histogram bin size

    nhis = (float*)malloc(sizeof(float)*NHBIN);
    //for (ih=0; ih<NHBIN; ih++) nhis[ih] = 0.0; // Reset the histogram

    cudaMalloc((void**)&dev_r,sizeof(float)*3*n);
    cudaMalloc((void**)&dev_nhis,sizeof(float)*NHBIN);

    cudaMemcpy(dev_r,r,3*n*sizeof(float),cudaMemcpyHostToDevice);

```

```

cudaMemset(dev_nhis,0.0,NHBIN*sizeof(float));

cudaMemcpyToSymbol(DALTH,alth,sizeof(float)*3,0,cudaMemcpyHostToDevice);
cudaMemcpyToSymbol(DN,&n,sizeof(int),0,cudaMemcpyHostToDevice);
cudaMemcpyToSymbol(DDRH,&drh,sizeof(float),0,cudaMemcpyHostToDevice);

dim3 numBlocks(8,8,1);
dim3 threads_per_block(16,16,1);
gpu_histogram_kernel<<<numBlocks,threads_per_block>>>(dev_r,dev_nhis);

cudaMemcpy(nhis,dev_nhis,NHBIN*sizeof(float),cudaMemcpyDeviceToHost);
cudaFree(dev_r);
cudaFree(dev_nhis);

density = n/(al[0]*al[1]*al[2]);
/* Print out the histogram */
fp = fopen("pdf.d","w");
for (ih=0; ih<NHBIN; ih++) {
    gr = nhis[ih]/(2*M_PI*pow((ih+0.5)*drh,2)*drh*density*n);
    fprintf(fp,"%e %e\n",(ih+0.5)*drh,gr);
}
fclose(fp);
free(nhis);
}

/*-----*/
int main() {
/*-----*/

    int i;
    float cpu1,cpu2;

    /* Read the atomic position data */
    fp = fopen("pos.d","r");
    fscanf(fp,"%f %f %f",&(al[0]),&(al[1]),&(al[2]));
    fscanf(fp,"%d",&n);
    r = (float*)malloc(sizeof(float)*3*n);
    for (i=0; i<n; i++)
        fscanf(fp,"%f %f %f",&(r[3*i]),&(r[3*i+1]),&(r[3*i+2]));
    fclose(fp);

    /* Compute the histogram */
    cpu1 = ((float) clock())/CLOCKS_PER_SEC;
    histogram();
    cpu2 = ((float) clock())/CLOCKS_PER_SEC;

```

```

printf("Execution time (s) = %le\n",cpu2-cpu1);

free(r);
return 0;
}

```

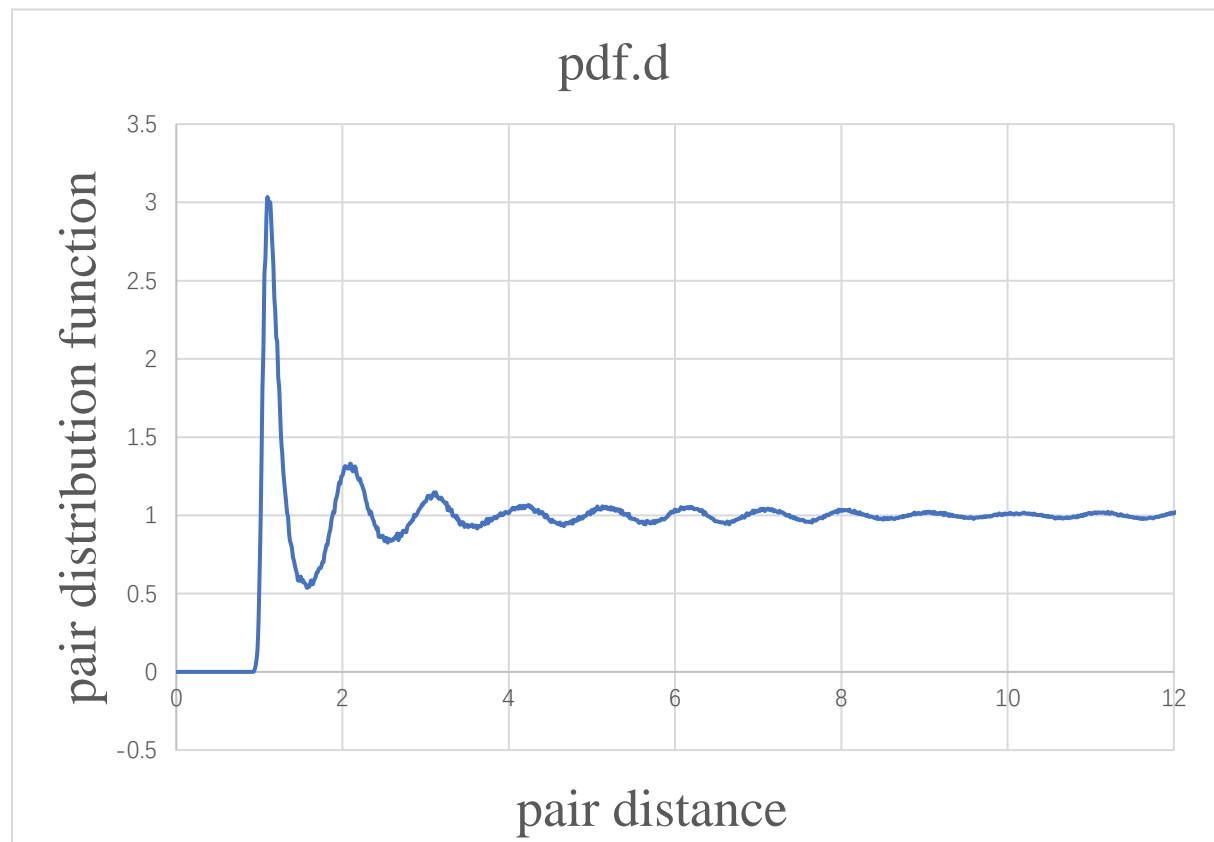
Output:

```

[youzhiqu@hpc-login3:~/work596$ more pdf.out
----- Begin SLURM Prolog -----
Job ID:      2265101
Username:    youzhiqu
Accountname: lc_an2
Name:        pdf.sl
Partition:   quick
Nodelist:    hpc3049
TasksPerNode: 1
CPUsPerTask: Default[1]
TMPDIR:      /tmp/2265101.quick
SCRATCHDIR:  /staging/scratch/2265101
Cluster:     uschpc
HSDA Account: false
----- 2018-11-16 20:52:12 -----
##### CPU: gcc -O -o pdf0 pdf0.c -lm #####
Execution time (s) = 2.660000e+00
##### GPU: nvcc -O -o pdf1 pdf1.cu #####
Execution time (s) = 2.300000e-01

```

2:



Part 2:

pi3.cu code:

```
// Hybrid MPI+OpenMP+CUDA computation of Pi
#include <stdio.h>
#include <mpi.h>
#include <omp.h>
#include <cuda.h>

#define NBIN 10000000 // Number of bins
#define NUM_DEVICE 2 // # of GPU devices = # of OpenMP threads
#define NUM_BLOCK 13 // Number of thread blocks
#define NUM_THREAD 192 // Number of threads per block

// Kernel that executes on the CUDA device
__global__ void cal_pi(float *sum,int nbin,float step,float offset,int nthreads,int nblocks) {
    int i;
    float x;
    int idx = blockIdx.x*blockDim.x+threadIdx.x; // Sequential thread index across the
    blocks
    for (i=idx; i<nbin; i+=nthreads*nblocks) { // Interleaved bin assignment to threads
        x = offset+(i+0.5)*step;
        sum[idx] += 4.0/(1.0+x*x);
    }
}

int main(int argc,char **argv) {
    int myid,nproc,nbin,tid;
    int mpid;
    float step,offset,pi=0.0,pig;
    dim3 dimGrid(NUM_BLOCK,1,1); // Grid dimensions (only use 1D)
    dim3 dimBlock(NUM_THREAD,1,1); // Block dimensions (only use 1D)
    float *sumHost,*sumDev; // Pointers to host & device arrays
    int dev_used;

    MPI_Init(&argc,&argv);
    MPI_Comm_rank(MPI_COMM_WORLD,&myid); // My MPI rank
    MPI_Comm_size(MPI_COMM_WORLD,&nproc); // Number of MPI processes
    //nbin = NBIN/nproc; // Number of bins per MPI process
    //step = 1.0/(float)(nbin*nproc); // Step size with redefined number of bins
    //offset = myid*step*nbin; // Quadrature-point offset

    omp_set_num_threads(NUM_DEVICE); // One OpenMP thread per GPU device
    nbin = NBIN/(nproc*NUM_DEVICE); // # of bins per OpenMP thread
    step = 1.0/(float)(nbin*nproc*NUM_DEVICE);
```

```

#pragma omp parallel private(mpid,offset,SumHost,sumDev,tid,dev_used)
reduction(+:pi)
{
    mpid = omp_get_thread_num();
    offset = (NUM_DEVICE*myid+mpid)*step*nbin; // Quadrature-point offset
    cudaSetDevice(mpid%2);
    //cudaSetDevice(myid%2);
    size_t size = NUM_BLOCK*NUM_THREAD*sizeof(float); //Array memory size
    sumHost = (float *)malloc(size); // Allocate array on host
    cudaMalloc((void **) &sumDev,size); // Allocate array on device
    cudaMemset(sumDev,0,size); // Reset array in device to 0
    // Calculate on device (call CUDA kernel)
    cal_pi <<<<dimGrid,dimBlock>>>>
(sumDev,nbin,step,offset,NUM_THREAD,NUM_BLOCK);
    // Retrieve result from device and store it in host array
    cudaMemcpy(sumHost,sumDev,size,cudaMemcpyDeviceToHost);
    // Reduction over CUDA threads
    for(tid=0; tid<NUM_THREAD*NUM_BLOCK; tid++)
        pi += sumHost[tid];
    pi *= step;
    // CUDA cleanup
    free(sumHost);
    cudaFree(sumDev);
    cudaGetDevice(&dev_used);
    //printf("myid = %d: device used = %d; partial pi = %f\n",myid,dev_used,pi);
    printf("myid = %d; mpid = %d: device used = %d; partial pi = %f\n", myid, mpid,
dev_used, pi);
} // End omp parallel

// Reduction over MPI processes
MPI_Allreduce(&pi,&pig,1,MPI_FLOAT,MPI_SUM,MPI_COMM_WORLD);
if (myid==0) printf("PI = %f\n",pig);

MPI_Finalize();
return 0;
}

```

2. output:

```

[youzhiqu@hpc-login3:~/work596$ salloc --nodes=2 --ntasks-per-node=2 --cpus-per-task=1 --gres=gpu:2 -t 29
salloc: Granted job allocation 2265134
salloc: Waiting for resource configuration
salloc: Nodes hpc[3049,3052] are ready for job
----- Begin SLURM Prolog -----
Job ID:      2265134
Username:    youzhiqu
Accountname: lc_an2
Name:        sh
Partition:   quick
Nodelist:    hpc[3049,3052]
TasksPerNode: 2(x2)
CPUsPerTask: 1
TMPDIR:      /tmp/2265134.quick
SCRATCHDIR:  /staging/scratch/2265134
Cluster:     uschpc
HSDA Account: false
----- 2018-11-16 21:10:29 -----
youzhiqu@hpc3049:/auto/dr-std/an2/youzhiqu$ source /usr/usc/openmpi/default/setup.sh
[youzhiqu@hpc3049:/auto/dr-std/an2/youzhiqu$ source /usr/usc/cuda/default/setup.sh
[youzhiqu@hpc3049:/auto/dr-std/an2/youzhiqu$ srun -n 2 ./pi3
myid = 1; mpid = 0: device used = 0; partial pi = 0.719409
myid = 1; mpid = 1: device used = 1; partial pi = 0.567582
myid = 0; mpid = 0: device used = 0; partial pi = 0.979926
myid = 0; mpid = 1: device used = 1; partial pi = 0.874671
PI = 3.141588
[youzhiqu@hpc-login3:~/work596$ salloc --nodes=2 --ntasks-per-node=2 --cpus-per-task=1 --gres=gpu:2 -t 29
salloc: Granted job allocation 2265134
salloc: Waiting for resource configuration
salloc: Nodes hpc[3049,3052] are ready for job
----- Begin SLURM Prolog -----
Job ID:      2265134
Username:    youzhiqu
Accountname: lc_an2
Name:        sh
Partition:   quick
Nodelist:    hpc[3049,3052]
TasksPerNode: 2(x2)
CPUsPerTask: 1
TMPDIR:      /tmp/2265134.quick
SCRATCHDIR:  /staging/scratch/2265134
Cluster:     uschpc
HSDA Account: false
----- 2018-11-16 21:10:29 -----
youzhiqu@hpc3049:/auto/dr-std/an2/youzhiqu$ source /usr/usc/openmpi/default/setup.sh
[youzhiqu@hpc3049:/auto/dr-std/an2/youzhiqu$ source /usr/usc/cuda/default/setup.sh
[youzhiqu@hpc3049:/auto/dr-std/an2/youzhiqu$ srun -n 2 ./pi3
myid = 1; mpid = 0: device used = 0; partial pi = 0.719409
myid = 1; mpid = 1: device used = 1; partial pi = 0.567582
myid = 0; mpid = 0: device used = 0; partial pi = 0.979926
myid = 0; mpid = 1: device used = 1; partial pi = 0.874671
PI = 3.141588

```